


Introduction into the Posix threads (Pthreads)


Role and place of parallel languages and libraries

* - по данным опроса

Suess M, Leopold C. Observations on the Publicity and Usage of Parallel Programming Systems and Languages: A Survey Approach. (2007)



The most popular parallel programming libraries *

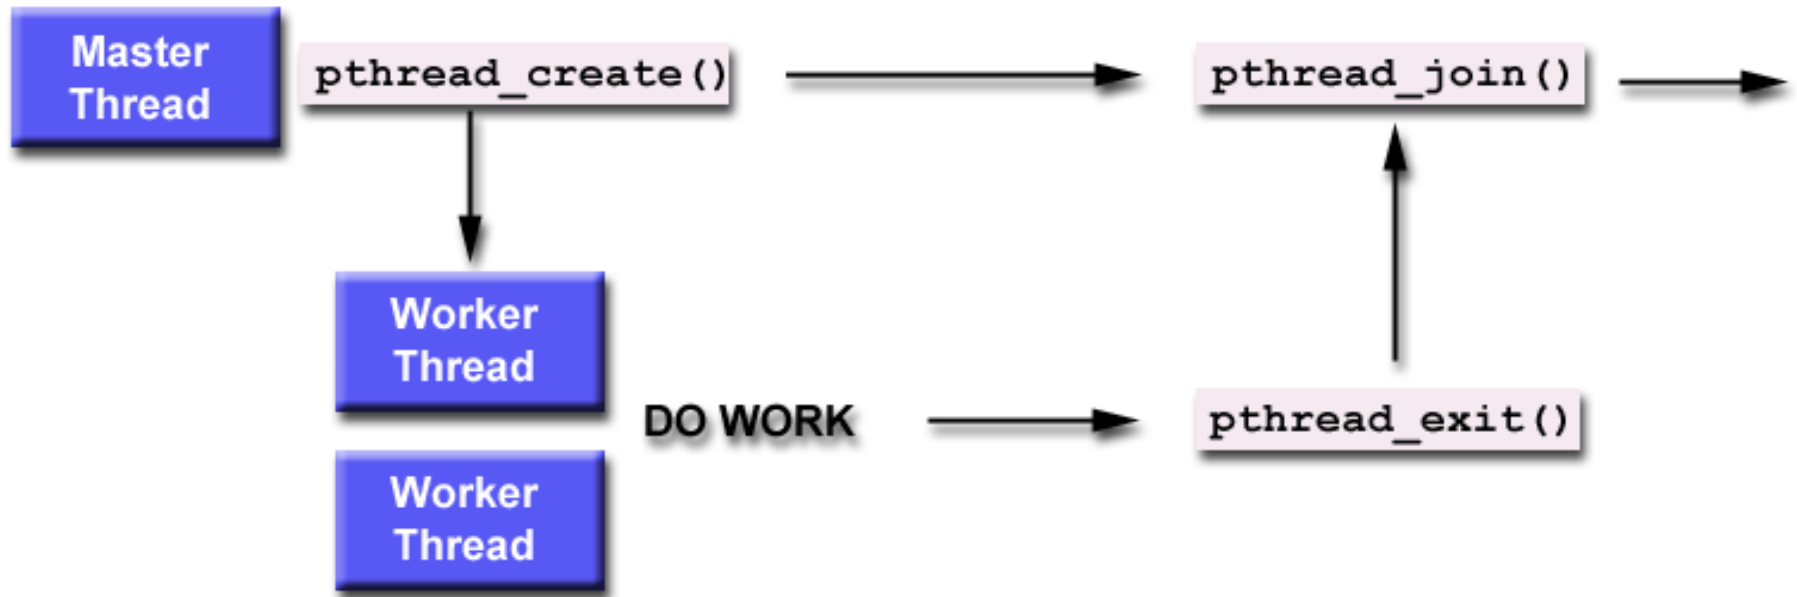


The most popular parallel programming languages *

Programming with Posix threads

- Thread is very similar to a process. It is also called as lightweight process.
- The main differences from a process. Process has its own memory area, independent of other processes, table of open files, current directory, and other kernel-level information. All threads, belonging to the same process, have all these entities in common.
- The Posix standard was adopted in 1995 and exists on UNIX and Linux-like systems (until 1995 it was only on UNIX-like systems). Almost all existing operating systems are compatible or practically compatible with this standard: BSD, Mac OS, Solaris, etc. There is an implementation of the corresponding libraries for Windows OS.
- In addition to Posix threads, there are the following common programming implementations for systems with shared memory (SMP - symmetric multiprocessing): OpenMP, Windows threads.

Scheme of a parallel program using Pthreads



Compilation and running

Compilation :

```
gcc    file_name.c  -lpthread  -lrt
```

(by default the executable is named "a.out")

Running:

```
./a.out  -  as usual
```

1st program, part 1

thread creation

```
#include<stdio.h>
```

```
#include<stdlib.h>
```

```
#include<pthread.h>    // pthread header file
```

```
int main(int argc, char *argv[]){
```

```
int param;
```

```
int rc;
```

```
void *arg;
```

```
pthread_t pthr;    // data type pthread
```

```
1) rc = pthread_create(&pthr, NULL, start_func, NULL); // thread  
creation, simplified call
```

```
2) rc = pthread_create(&pthr, NULL, start_func, (void*) &param);
```

1st program, part 2

/* thread finishes when it exits the start_func function. If you need to get the return value of a function, then you need to use */

- 1) `pthread_join(pthr, NULL);` // accept nothing by default
 - 2) `pthread_join(pthr, &arg);` // write the address of the returned variable into &arg
- } // end of the main function

1st program, part 3

thread's function

```
/* Thread's work function (thread's life area) */
```

```
void* start_func(void* param){
```

```
    int *local;
```

```
    local = (int*) param;
```

```
    1) pthread_exit( NULL);    // used by default
```

```
    2) pthread_exit( (void*) &arg);    // it is used if it is necessary to  
terminate the execution of a thread earlier with a return value. It is  
possible to use NULL as an argument, the arg variable must be  
dynamically allocated or have a static memory class
```

```
    3) return NULL;    // or like this
```

```
}    // at the moment of exiting the start_func function, thread releases its  
resources (dies)
```


1st program, part 4

semaphores

Thread synchronization mechanism. There are several mechanisms for organizing access to the critical section.

Critical section is a piece of program code that contains a shared resource and is only accessible to one thread.

An N-dimensional semaphore takes values from 0 to N, we will only use a binary semaphore that takes values 0 and 1

1st program, part 5

semaphores

```
#include<semaphore.h>
```

```
sem_t sem; // semaphore declaration – as a global variable
```

```
int main(int argc, char *argv[]){
```

```
    sem_init (&sem, 0, 1); // semaphore initialization
```

attribute



initial value



```
...
```

```
    sem_destroy(&sem); // releasing of the semaphore
```

```
    return 0;
```

```
}
```

1st program, part 6

semaphores

```
void* start_func(void* param){  
    int val;  
  
    ...  
  
    {  
        sem_wait(&sem); // decreases the semaphore value by 1  
                        // if sem = 0 at the time of this function  
                        // execution, then thread waits ...  
                        // actions on a shared variable  
        sem_post(&sem); // increases the semaphore value by 1  
  
        sem_getvalue(&sem, &val); // checking the semaphore value  
  
        return NULL;  
    }  
}
```

Critical section

Some software design

// several threads creation

```
#define NUM_THREADS 2
int main (int argc, char *argv[])
{
    pthread_t pthr[NUM_THREADS];
    void *arg;
    ...
    for(i = 0; i < NUM_THREADS; i++){
        rc = pthread_create(&pthr[i], NULL, start_func, NULL);
        if (rc) printf("ERROR; return code from pthread_create() is %d \n", rc);
    }
    ...
    for(i = 0; i < NUM_THREADS; i++){
        rc = pthread_join(pthr[i], &arg);
        printf("value from func %d \n", *(int*)arg);    // example
        if (rc) printf("ERROR; return code from pthread_join() is %d \n", rc);
    }
}
```

Determining the running time of a parallel program

```
#include<time.h> // header file containing types and functions for working  
with date and time
```

```
struct timespec begin, end; // requires key -lrt !  
double elapsed;
```

```
clock_gettime(CLOCK_REALTIME, &begin); /* returns a reference to a record  
of type timespec, which is declared in  
time.h and has fields:  
time_t tv_sec; – seconds,  
long tv_nsec; – nanoseconds.  
*/
```

```
... // threads work here
```

```
clock_gettime(CLOCK_REALTIME, &end);
```

```
elapsed = end.tv_sec - begin.tv_sec; // time in seconds  
elapsed += (end.tv_nsec - begin.tv_nsec) / 1000000000.0; // add time down  
to nanoseconds
```

On the function of generating random numbers 1

The rand () function is not thread-safe. It looks like this:

```
unsigned int next = 1;
```

```
/* rand - return pseudo-random integer on 0..32767 */
```

```
int rand(void) {  
    next = next*1103515245 + 12345;  
    return (unsigned int)(next/65536) % 32768;  
}
```

```
/* srand - set seed for rand() */
```

```
void srand(unsigned int seed) {  
    next = seed;  
}
```

and deals with the global intermediate variable next, which changes every time each thread is called. This intermediate value is stored in a static memory area, that is, it is shared. In addition to the non-repeatability of the result, this leads to a strong slowdown in the program.

On the function of generating random numbers 2

An analogue of the `rand ()` function is supplied with the Posix standard library, but already thread-safe. Its appearance:

```
/* rand_r – a reentrant pseudo-random integer on 0..32767 */
int rand_r(unsigned int *nextp) {
    *nextp = *nextp * 1103515245 + 12345;
    return (unsigned int)(*nextp / 65536) % 32768;
}
```

that is, the function deals with a thread local variable to be declared in the thread handler function. Such a function is called reentrant. Since the intermediate value is already a local variable, there is no access conflict.

In a program:

```
int x_k, y_k, z_k;

x = ((float)rand_r(&x_k) / RAND_MAX) * (i_final - i_init) + i_init;
y = (float)rand_r(&y_k) / RAND_MAX;
z = (float)rand_r(&z_k) * Z0 / RAND_MAX;
```